

# Lab 7

Gretel Warmuth (PID: A17595945)

Today we are going to learn how to apply different machine learning methods, beginning with clustering:

The goal is to find groups/clusters in the input data.

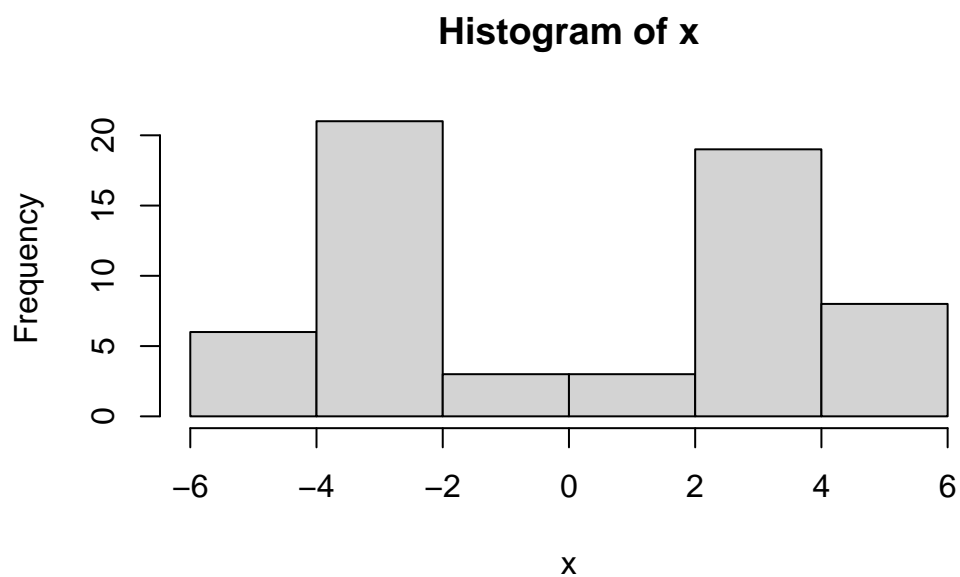
First I will make some random data with clear groups with the `rnorm()` function:

```
hist(rnorm(1000, mean=3))
```



Making a histogram with two peaks:

```
n <- 30  
x <- (c(rnorm(n, -3), rnorm(n, +3)))  
hist(x)
```



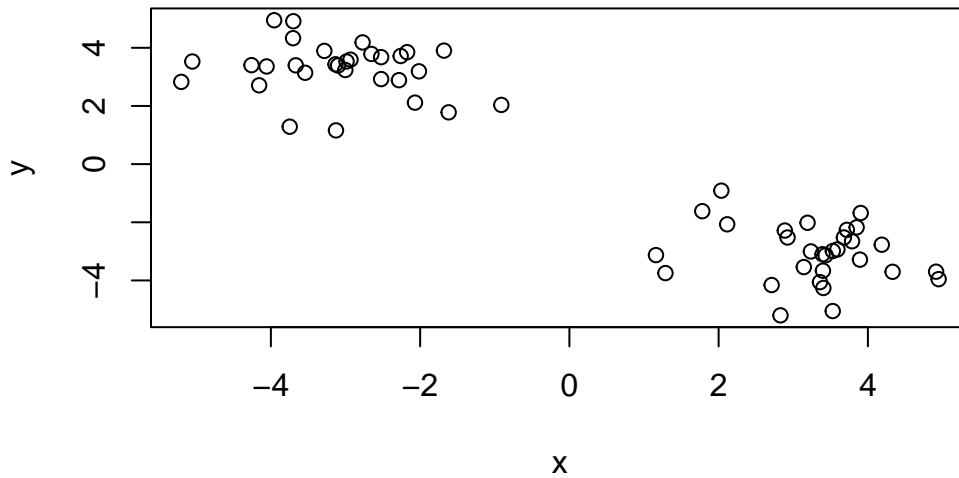
Making a cluster plot:

```
n <- 30
x <- (c(rnorm(n, -3), rnorm(n, +3)))
y <- rev(x)

z<- cbind(x,y)
head(z)
```

```
      x      y
[1,] -1.617962 1.780855
[2,] -5.052637 3.528754
[3,] -2.522329 2.923445
[4,] -3.002405 3.236163
[5,] -3.699434 4.913742
[6,] -2.770901 4.185858
```

```
plot(z)
```



Use the `kmeans()` function setting `k` to 2 and `nstart=20`

Inspect/print the results

Q. How many points are in each cluster?

Q. What ‘component’ of your result object details - cluster size? - cluster assignment/membership? - cluster center?

Q. Plot  $x$  colored by the kmeans cluster assignment and add cluster centers as blue points

```
km <- kmeans(z, centers = 2)
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	-3.069540	3.271348
2	3.271348	-3.069540

Clustering vector:

[illegible]

```
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

Within cluster sum of squares by cluster:

```
[1] 52.31094 52.31094
(between_SS / total_SS = 92.0 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

Results of the object km

```
attributes(km)
```

```
$names
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

```
$class
[1] "kmeans"
```

Cluster size:

```
km$size
```

```
[1] 30 30
```

Cluster assignment/membership:

```
km$cluster
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

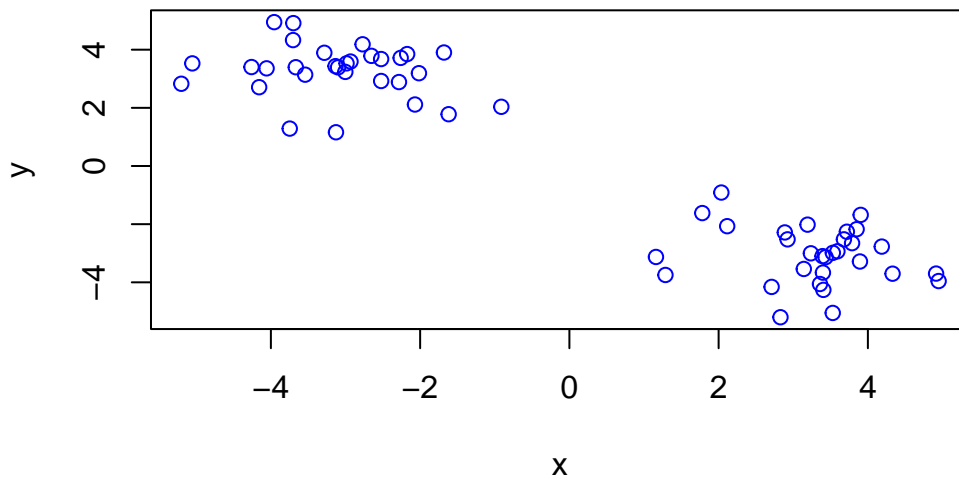
Cluster center:

```
km$centers
```

	x	y
1	-3.069540	3.271348
2	3.271348	-3.069540

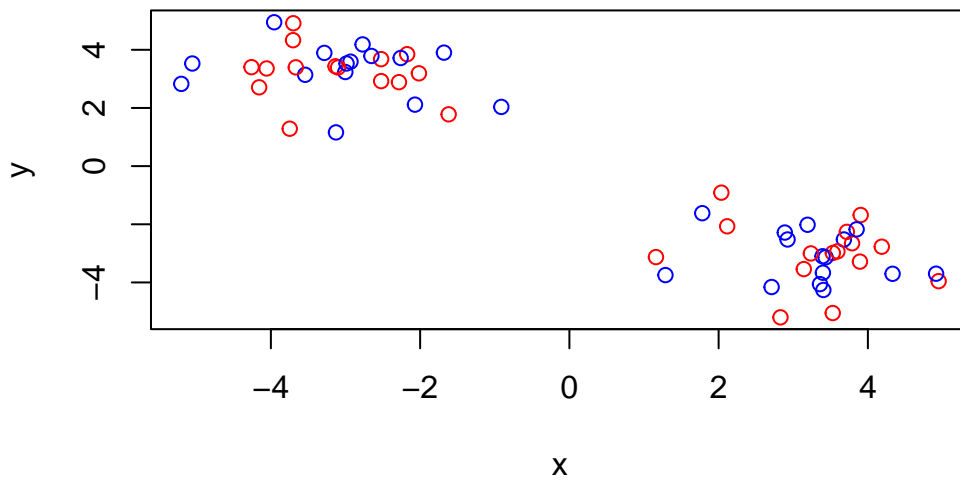
Plot:

```
plot(z, col="blue")
```



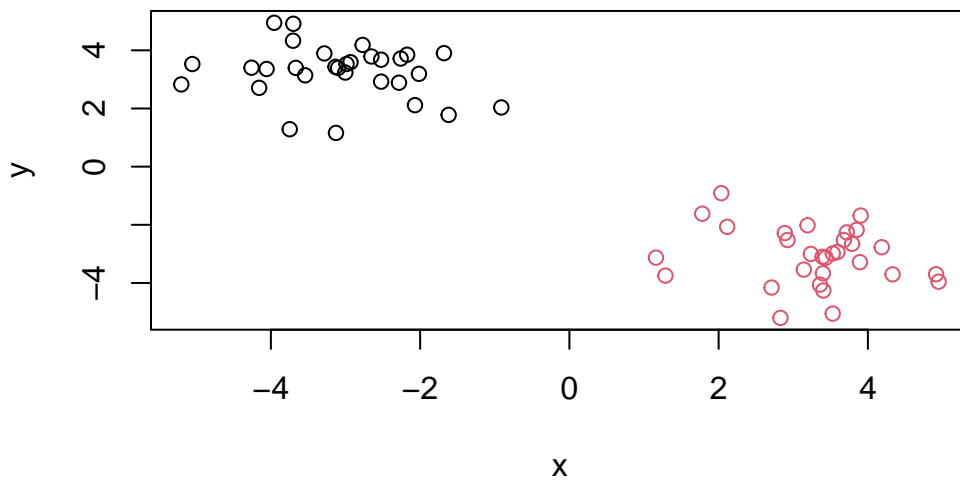
R will recycle the shorter color vector to be the same length as the longer (number of data points) in `z`

```
plot(z, col=c("red", "blue"))
```



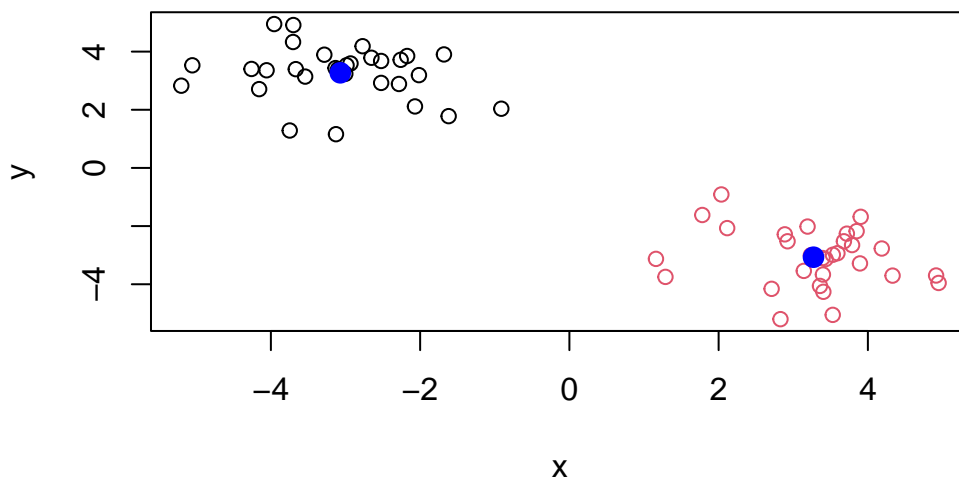
Coloring the clusters:

```
plot(z, col=km$cluster)
```



We can use the `points()` function to add new points to an existing plot like for the cluster centers:

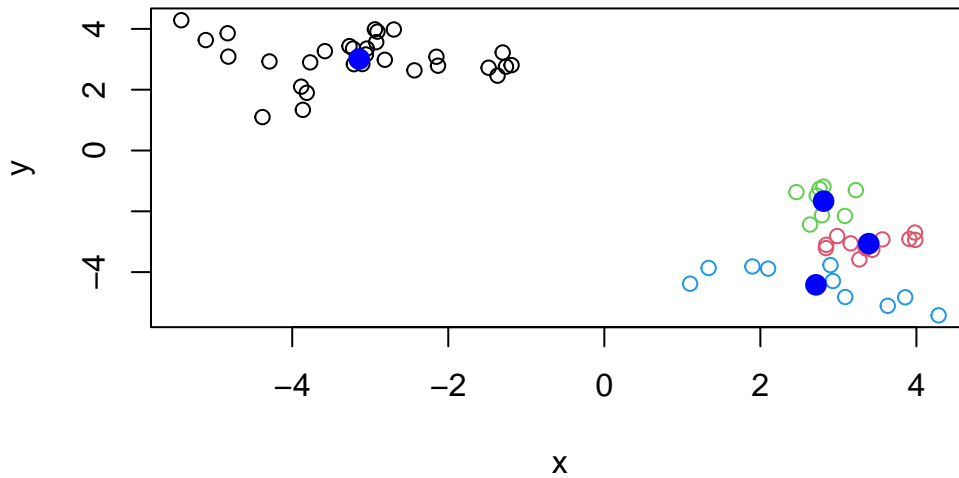
```
plot(z, col=km$cluster)
points(km$centers, col="blue", pch=16, cex=1.5)
```



Q. Run km again and ask for 4 clusters and plot them

```
n <- 30
x <- (c(rnorm(n, -3), rnorm(n, +3)))
y <- rev(x)

z<- cbind(x,y)
km4 <- kmeans(z, centers = 4)
plot(z, col = km4$cluster)
points(km4$centers, col="blue", pch=16, cex=1.5)
```



## Hierarchical Clustering

Let's take our same made-up data `z` and see how `hclust` works.

First we make a distance matrix of our data to be clustered:

```
d <- dist(z)
hc <- hclust(d)
hc
```

Call:

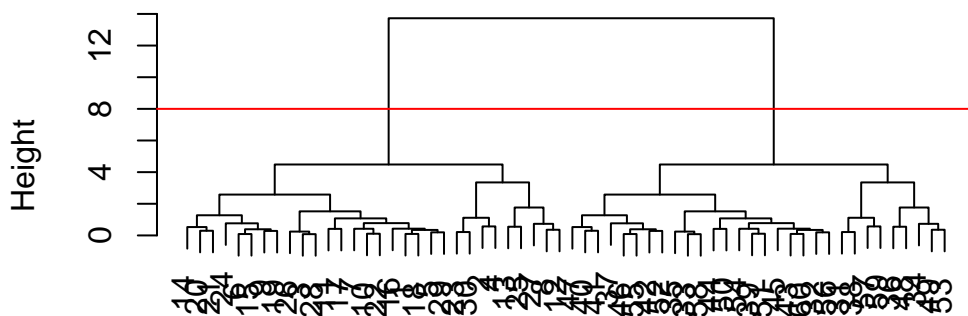
```
hclust(d = d)
```

```
Cluster method : complete
Distance       : euclidean
Number of objects: 60
```

```
plot(hc)
abline(h=8, col="red")
```



## Cluster Dendrogram



```
hclust (*, "complete")
```

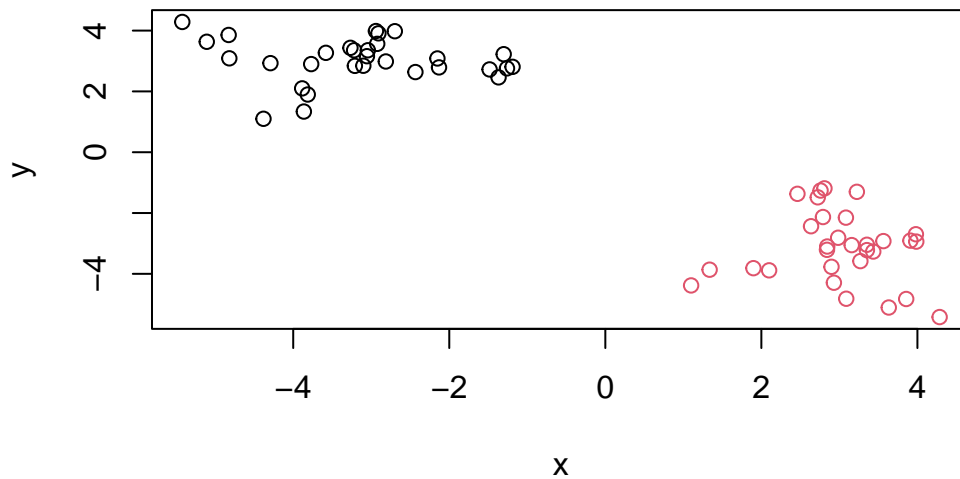
I can get my cluster membership vector by “cutting the tree” with the `cutree()` function:

```
grps <- cutree(hc, h=8)
grps
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

Can you plot `z` colored by our hclust results:

```
plot(z, col=grps)
```



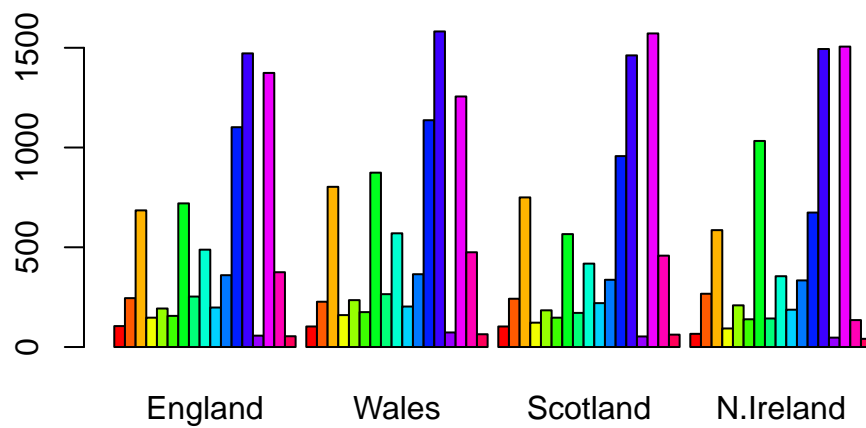
## PCA of UK Food Data

Read data from the UK on food consumption in different areas

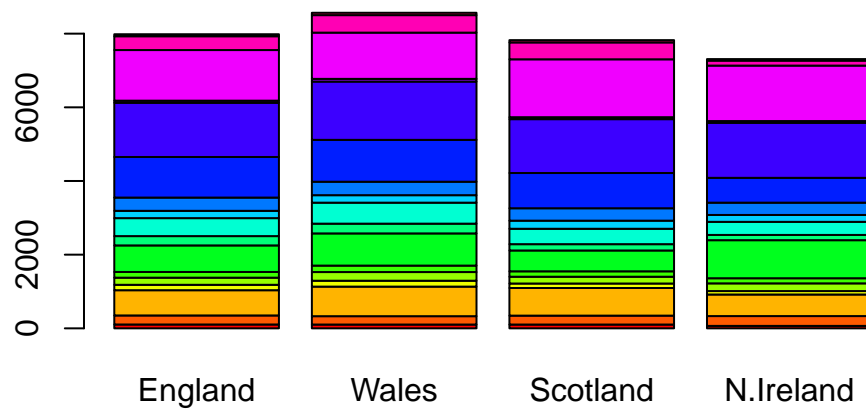
```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names = 1)
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```

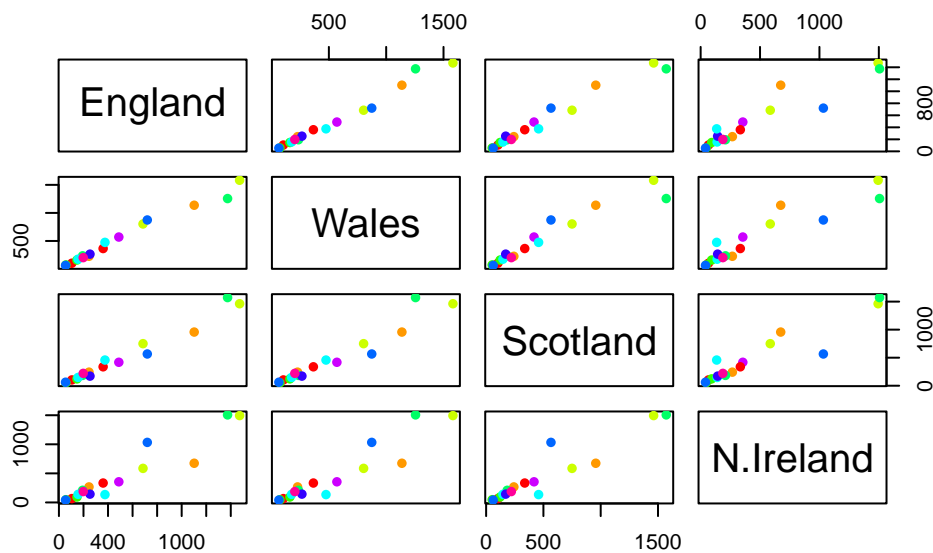


```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



A so-called pair-wise plot may help to compare countries and categories

```
pairs(x, col=rainbow(10), pch=16)
```



It is difficult to see structure and trends in this small dataset- how can we compare when we have even larger data?! PCA to the rescue!

## PCA

The main function in base R to do PCA is called `prcomp()`

```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	3.176e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Let's look at the `pca` object that we created from running `prcomp()`

```
attributes(pca)
```

```
$names  
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

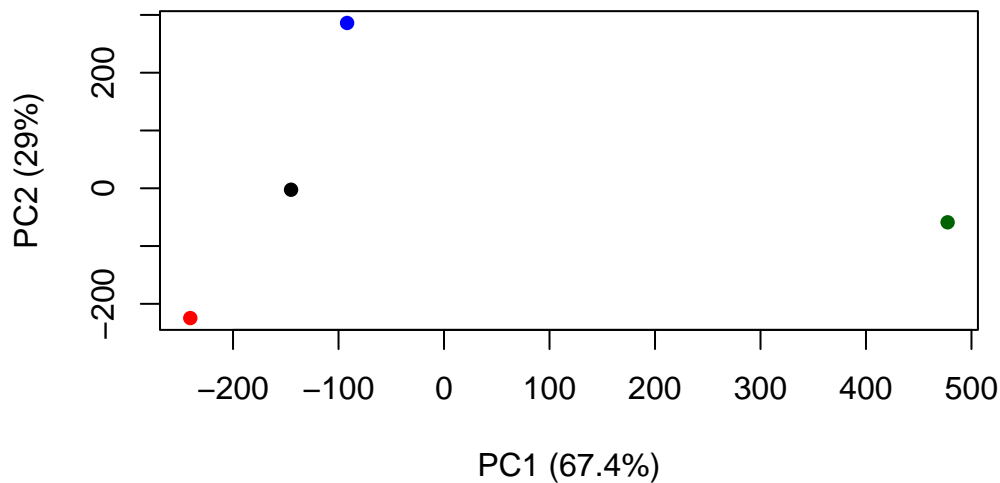
```
$class  
[1] "prcomp"
```

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-4.894696e-14
Wales	-240.52915	-224.646925	-56.475555	5.700024e-13
Scotland	-91.86934	286.081786	-44.415495	-7.460785e-13
N.Ireland	477.39164	-58.901862	-4.877895	2.321303e-13

The PCA plot:

```
plot(pca$x[,1], pca$x[,2],  
     col=c("black", "red", "blue", "darkgreen"), pch=16,  
     xlab="PC1 (67.4%)", ylab="PC2 (29%)")
```



Bar plot:

```
par(mar=c(10, 3, 0.35, 0))  
barplot( pca$rotation[,1], las=2 )
```

